

#pcodes = 2^4=16  
 #pos = 2^12=4096  
 #sizeOfPos = 16bits=2bytes  
 #sizeOfMemory=#pos \* sizeOfPos  
 #memoryRange = 000-FFF  
 #DataRange=0000-FFFF  
 #PC=3octalDigits=log2(8)=3bits\*3=9 bits  
 #PC is length of mem address!!!  
 #IR = opcode + address  
 #IR=Data+memPos=12bits  
 #pos = 2^9 = 512  
 #sizeOfAPosition=12bits  
 #SizeOfMemory= #pos \* #sizeOfAPosition  
 #memoryRange=000-777  
 #DataRange=0000-7777

## Process Control (continued)

**Modes of Execution:** \*User mode (less-privileged mode) -user programs typically execute in this mode. \*System mode (more privileged) - also referred to as control mode or kernel mode - kernel of the OS.

**KERNEL functions** - \*Process management (process creation, termination, switching, synchronization, management of proc. control block). \*Memory management (allocation of address space to process, swapping, page/segment management). \*I/O management (buffer management, Allocation of I/O channels and devices to process). \*Support functions (interrupt handling, Accounting, Monitoring).

**Process creation (step-by-step):** 1) OS assigns a unique process ID to the new process 2) allocates space for the process. 3) Initializes the process control block. 4) sets the appropriate linkages 5) creates or expands other data structures.

**Process Switching** - a process switch may occur anytime that the OS has gained control from the currently running process. Events giving OS control are \*interrupt (reaction an asynchronous external event)

\*Trap (Holding of an error or an exception condition) \*Supervisor call (call to an operating system function)

**System Interrupt:** \*due to some sort of event that is external to and independent of the currently running process. \*clock interrupt \*I/O interrupt \*memory fault \*time slice - the max amount of time that a process can execute before being interrupted.

**Trap:** \*an error or exception condition generated within the currently running process \*OS determines if the condition is fatal - moved to the exit state and a process switch occurs \*action will depend on the nature of the error.

**If no interrupts are pending the Processor:** 1) proceeds to the fetch stage and fetches the next instruction of the current program in the current process.

**If interrupt is pending:** 1) sets the program counter to the starting address of an interrupt handler program. 2) switch from user mode to kernel mode so that the interrupt processing code may include privileged instructions.

**Change Of Processor State:** 1) save the context of the proc 2) update the process control block of the process current running. 3) move the process control block of this process to the appropriate queue. 4) select another process for exe. 5) update the process control block of the process selected. 6) update memory management data structures. 7) restore the context of the processor to that which existed at the time the selected process was last switched out.

**Security Issues:** \* An OS associates a set of privileges with each process \*typically a process that exe on behalf of a user has the privileges that the OS recognizes for that user. \*Highest lvl of privilege is referred to as admin, supervisor, root.

**System Access Threats:** \*Intruders (hacker) \*Malicious Software. Intrusion Detection counter measure: Intrusion detection system (IDS\_ comprises of three loical components 1)sensors 2)analizers 3)user interface \*IDS are designed to detect human intruder behavior.

**Authentication Countermeasure:** consists of two steps 1) Identification 2) verification

**Access Control Countermeasure:** \*implements a security policy that specifies who or what may have access to each specific system resource and the type of access that is permitted in each instance.

\*mediates between a user and system resource. \*A security admin maintains an authorization database. \*auditing function monitors and keeps a record of user accesses to system resources.

**FIREWALL:** dedicated computer that: \*interfaces with comps outside the network. \*has special security precautions built into it to protect sensitive files on computers within the network. \*\*Design of a firewall: \*all traffic must pass through the firewall. \* only authorized traffic will be allowed to pass. \*immune to penetration.

**Fork() process creation:** 1) allocate a slot in the process table for the new process. 2) assign a unique process ID to the child process. 3) Make a copy of the process image of the parent, with the exception of any shared memory. 4) increments counters for any files owned by the parent, to reflect that an additional process now also owns those files. 5) Assings the child process to the ready to run state. 6) Returns the ID number of the child to the parent process, and a 0 value to the child process.

**After fork() process creation:** the kernel can do: 1) stay in the parent process. 2) An OS controls to the child process. 3) transfer control to another process.




Figure 3.36

Process Image: Operating System Executes within User Space

```

struct operation *pos_ptr = (struct operation *)pos_void_ptr;

if(pthread_create(&tid[i], NULL, calculator, &operations[i]))


for (int i = 0; i < NOPER; i++)
pthread_join(tid[i], NULL);
  
```

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

Table 4.2 Relationship between Threads and Processes

```

#include<iostream>
#include<unistd.h>
#include<sys/wait.h>
using namespace std;
int main()
{
    pid_t pid;
    int i = 0;
    for(i;i<3;i++){
        pid=fork();
        if(pid==0){
            break;
        }
        wait(NULL);
    }
    if(i == 0 && pid == 0){
        for(int j=0;j<2;j++){
            pid=fork();
            if(pid == 0){
                break;
            }
            wait(NULL);
        }
    }
    if(i == 2 && pid == 0){
        for(int j = 0; j<2; j++){
            pid = fork();
            if(pid == 0){
                break;
            }
            wait(NULL);
        }
    }
}
  
```



(a) With One Suspend State





Figure 3.6 Five-State Process Model



```

#include<iostream>
#include<unistd.h>
#include<sys/types.h>
#include<sys/wait.h>
int main()
{
    pid_t pid;
    std::cout << "I am the parent process" << std::endl;
    for(int i=0;i<3;i++){
        pid = fork();
        if (pid == 0)
        {
            std::cout << "I am the child process " << i << std::endl;
            if (i==1)
            {
                pid = fork();
                if (pid == 0)
                {
                    std::cout << "I am a grandchild process from child process " << i << std::endl;
                    _exit(0);
                }
                wait(0);
            }
            _exit(0);
        }
        return 0;
    }
}
  
```

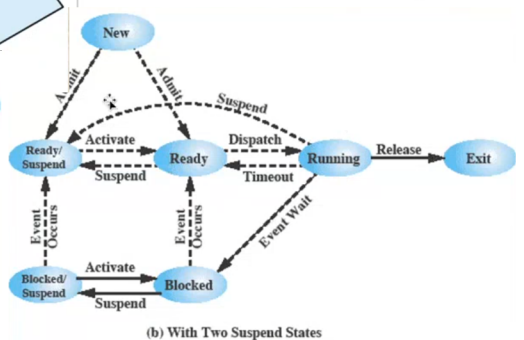
# Blocks in Main Memory =  $2^n / K$

**Two levels of memory(cache/main)**  
**Hit Ratio:** The probability of a word being found in the first level of memory (fastest)  
**Miss Ratio:** 1 - Hit ration

**Principle of Locality (cache)-** when you transfer information from second level to first level, you will transfer a block of words where that block is. This is because the probability of another word we are looking for is higher. Code tends to be sequential. Increases performance

## PERFORMANCE OF A SIMPLE TWO-LEVEL MEMORY

Access time Level 1 (TL<sub>1</sub>) = 100 ms  
 Access time Level 2 (TL<sub>2</sub>) = 1000 ms  
 Hit Ratio = 90 %  
 Miss Ratio = 10%  
 Average Access time (AvgT) = HR \* TL<sub>1</sub> + MR \* (TL<sub>2</sub>+TL<sub>1</sub>)  
 = 0.9 \* 100 ms + 0.1 \* (1000 ms + 100 ms)



(b) With Two Suspend States




Figure 4.2 Single Threaded and Multithreaded Process Models

**Benefits of Threads:** 1) takes less time to create a new thread vs. new process. 2) less time to terminate a thread than a process. 3) switching between threads takes less time than switching between processes. 4) threads enhance efficiency in communication between programs.

**Thread use in a single-user system:** \*foreground and background work, \*asynchronous processing, \*speed of execution, \*modular program structure.

**Suspending a process/thread:** \*involves suspending all threads of the process.

**termination of a process terminates all threads within the process.**

**Thread execution state:** \*key states (running, ready, blocked). \*Thread operations associated with a change in thread state are: 1) spawn 2) block 3) unblock 4) finish.

**thread synchronization:** \*It is necessary to synchronize the activities of the various threads. \*all threads of a process share the same address space and other resources. \*any alternation of a resource by one thread affects the other threads in the same process.

**Types of threads:** 1) Upper level thread(ULT) 2)Kernel level thread(KLT)

**ULT Upper level threads:** \*all thread management is done by the application. \*the kernel is not aware of the existence of threads.

**ULT Advantages:** \*thread switching does not require kernel mode privileges, \*scheduling can be application specific, \*ULTs can run on any OS.

**ULT Disadvantages:** \*in a typical OS many system calls are blocking -as a result, when a ULT executes a system call, not only is that thread blocked, but all of the threads within the process are blocked.

\*In pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing.

**Overcoming ULT disadvantages:** \*Jacketing -converts a blocking system call into a non-blocking system call. \*writing an application as multiple processes rather than multiple threads.


**Kernel-Level Threads(KLTs):** \*Thread management is done by the kernel. \*no thread management is done by the application, \*Windows is an example of this approach.

**KLT Advantages:** \*The kernel can simultaneously schedule multiple threads from the same process on multiple processors. \*If one thread in a process is blocked, the kernel can schedule another thread of the same process. \*Kernel routines can be multithreaded.

**KLT Disadvantages:** The transfer of control from one thread to another within the same process requires a mode switch to the kernel.

**Combined Approaches:** \*Thread creation is done in the user space. \*Bulk of scheduling and synchronization of threads is by the application, \*Solaris is an example.

**Applications that Benefit:** \*Multithreaded native applications -characterized by having a small number of highly threaded processes. \*Multiprocess applications -characterized by the presence of many single-threaded process. \*Java Applications, \*Multiinstance applications - multiple instances of the application in parallel.



\* Operating system - exploits hardware resources of one or more processors, most complex pieces of software  
- provides a set of services to system users, manages second. mem and I/O dev.  
\* pc - holds address of the instruction to be fetched next (incr. after each fetch)  
PROCESSOR  
IR - Instruction being executed.  
MAR -  
MBR -  
I/O AR -  
I/O BR -  
\* AC (accumulator) - temporary storage  
microprocessor - fastest general purpose proc. proc on single chip, multiprocessor, each chip cont. cores  
GPU - provide efficient computation or arrays, physics simulation, comp. on large spreadsheet.  
\* basic instr cycle - start -> fetch inst -> execute inst. (repeat if needed) -> halt  
\* Interrupts - interrupt the normal sequencing of a proc  
provided to improve proc utilization (i/o devices are slower than the proc)  
\* classes of interrupt (Program, Timer, I/o, hardware failure)  
\* single interrupt process - dev. controller/sys harw issues inter -> proc finishes  
current instruc -> proc ack. interrupt -> proc pushes PSW and PC onto cntrl stack ->  
proc loads new PC val from inter -> save remainder of proc state info -> process interrupt ->  
restore proc state info -> restore old PSW and PC  
\* Multiple interrupts - can be done sequentially or nested  
\* MEMORY HIERARCHY major constraints (amount, speed, expense), memory must keep up with proc,  
cost of mem must be reasonable to other components  
\* down pyramid (decreasing cost per bit, increasing cap., increasing access time, decr. freq. of acc. to mem by proc.)  
\* PRINCIPLE OF LOCALITY - 1) memory references by the proc. tend to cluster 2) data is org. so that the percentage of  
accesses to each successively lower level is less than that of the lvl above 3) can be applied to more than 2 lvl's  
\* CACHE 1) invisible to OS. 2) interacts with mem mang. hardware. 3) Proc. must access mem. at least once per instr.  
cycle. 4) exploits principle of locality with small, fast mem.  
\* Cache Princ. 1) contains copy of portion of main mem. 2) Proc first checks cache 3) not found? block of mem is read  
into cache. 4) locality of ref- it is likely that many of future mem. references will be other bytes in block.  
MAPPING FUNCTION - Determines which cache location the block will occupy  
\* Replacement Algorithm - Least recently used (LRU) Algo (effective strat to replace a block that has been in the cache  
the longest with no refer. \* hardware mechanisms are needed to identify the least recently used block \* chooses which  
block to replace when a new block is to be loaded into cache)  
\* I/O Techniques - When the proc encounters an instr. relating to I/O, it exe that instr. by issuing a comm to the  
appropriate I/O module  
\* PROGRAMMED I/O - 1) the I/O module performs the requested action then sets the appropriate bits in the I/O  
status reg. 2) The Proc periodically checks the status of the I/O module until it determines the instruc. is complete.  
3) With programmed I/O the performance level of the entire systm is severely degraded  
Interrupt Driven I/O drawbacks - transfer rate is limited by the speed with which the proc can test/service a device.  
2) The proc is tied up in managing an I/O transf. - number of instruc. must be exe for each I/O transfer.  
DIRECT MEMORY ACCESS (DMA) - performed by a separate module on the sys bus or incorporated into an I/O module  
1) transfers the entire block of data directly to and from mem without going through the proc  
- proc is involved only at the beg. and end of transfer  
- proc exe more slowly during a transfer when proc access to the bus is req.

2) more efficient than interrupt-driven I/O  
Symmetric Multi-Proc (SMP) - stand-alone computer systm with  
- two or more similar processors of comparable capabil.  
- processors share the same main mem and are interconnected by bus  
- proc share access to I/O devices  
- all processors can perform the same functions  
- system is controlled by intergrated  
OS that provides interaction between processors and their programs  
at the job, task, file, data element levels.  
Multi-core comp (chip multi-processor) - combines two or more proc(cores) on a single silicon  
- each core consists of all components of an indepen. proc.  
- multicore chips include L2/L3 cache.

OPERATING SYSTEM - program that controls the execution of application programs, interface between application and hardware.  
\* Services OS - program devl., program exe, access i/o devices, controlled access to files, system access, error detection and response.  
Key Elements \* Instruction set arch. (ISA), application binary interf. (ABI)  
Application programming interface (API)  
Role of OS - computer is a set of resources for the movement, storage, and proc. of data, the OS is responsible for managing these resources.  
Operating system as software 1) functions in teh same way as ordinary comp software.  
2) program, or suite of programs, executed by the proc. 3) frequently relinquishes control and must depend on the proc. to allow it to regain control.  
Serial processing - (earliest computers, no OS, programs interacted directly with hardware, ran with display lights, toggle switches, printer, users had to interact with the computer in 'series') PROBLEMS: scheduling - time allocations could run short or long, wasted computer time. set up-time - a considerable amount of time was spent on just setting up the program to run.  
Simple Batch system - early computers were expensive, important to maximize proc utilization. \* Monitor - user no longer has direct access to processor,, job is submitting to computer operator who batches them together and palces them on an input device. Program branches back to the monitor when finished.  
Monitor Point of view - monitor control the sequence of events,, \* resident monitor software always in memory, \* monitor reads in jobs and gives control, \* job returns control to monitor.  
Processor POV - \* Proc exe instructions from the mem containing the monitor, \* Exe the instructions in the user program until it encounters an ending or ERR condition, \*\* Control is passed to a job" means proc is fetching and exe instructions in a user program, \*\* control is returned to the monitor" means that the proc is fetching and exe instructions from the monitor program.  
Simple Batch System Overhead - \* Proc alternates between execution of user program and execution of the monitor.  
\* Sacrifices: 1) some main mem is now given over to the monitor, 2) some proc time is consumed by the monitor.  
\* Despite overhead, the simple batch system improves utilization of the comp.  
Uniprogramming - The proc spends a certain amount of time exe, until it reaches an I/O instruction; it must then wait until that I/O instruction concludes before proceeding.  
Multiprogramming - \* there must be enough memory to hold the OS (resident monitor) and one user program, \* When one job needs to wait for I/O, the processor cans witch to the other job, which is likely not waiting for I/O  
\* also known as multitasking, \* mem is expanded to hold three, four, or more programs and switch among all of them.  
Time-sharing system - \* can be used to handle multiple interactive jobs, \* Proc time is shared among multiple users. \* Mult. users simultaneously access the system through terminals, with the OS interleaving the exe of each user program in a short burst or quantum of computation.  
Process - fundamental to the structure of operating systems.  
Development of the process - Three major lines of computer system dev. created problems in timing and synchronization that contributed to the dev.  
Causes of error: 1) improper synchronization - a program must wait until the data are available in a buffer, improper design of the signaling mechanism can result in loss or duplication.  
2) failed mutual exclusion - more than one user or program attempts to make user of a shared resource at the same time. \* only one routine at a time allowed to perform an update against the file.  
3) nondeterminate program operation - program exe is interleaved by the proc when mem is shared. \* the order in which programs are scheduled may affect their outcome  
4) Deadlocks - it is possible for two or more programs to be hung up waiting for each other. \* may depend on the chance timing of resource allocation and release.  
Components of a process: \* a process contains three components 1) an executable program 2) the associated data needed by the program (variables, work space, buffers) 3) The execution context "Process State" of the program  
The execution context: \* is the internal data by which the OS is able to supervise and control the process, \* includes the contents of the various process registers, \* includes information such as the priority of the process and whether the process is waiting for the completion of a particular I/O event.  
Process management: \* The entire state of the process at any instant is contained in its context, \* New features can be designed and incorporated into the OS by expanding the context to include any new information needed to support the feature.  
Memory Management: the OS has five principle storage management responsibilities  
Virtual memory: \* a facility that allows programs to address memory from logical point of view, without regard to the amount of main memory physically available. \* conceived to meet the requirement of having multiple user jobs reside in main memory concurrently.  
Paging: \* Allows processes to be comprised of a number of fixed-size blocks. \* Program references a word by means of virtual address (consists of a page number and an offset with the page), (each page my be located anywhere in main mem.). \* Provides for a dynamic mapping between the virtual address used in the program and a real address in main mem.




Figure 1.7 Interactive Cycle with Interrupts

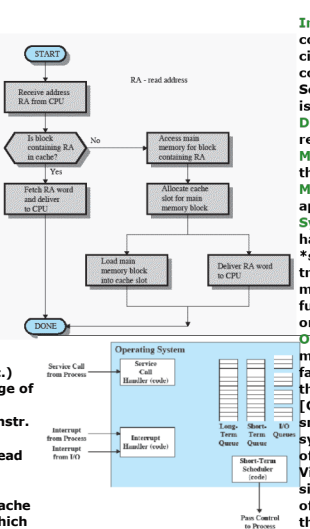
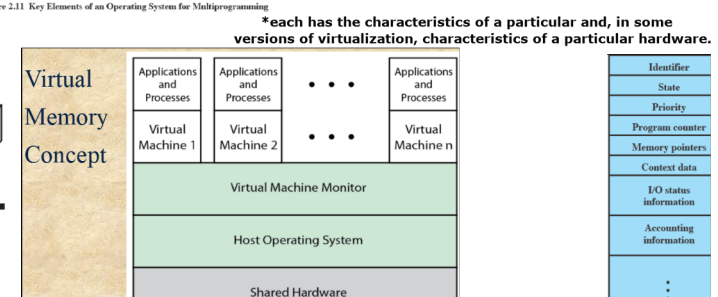


Figure 2.11 Key Elements of an Operating System for Multiprogramming

Information Protection and Security: the nature of the threat that concerns an organization will vary greatly depending on the circumstances. \*The problem involves controlling access to computer systems and the information stored in them.  
Scheduling and Resource Management: \*key responsibility of an OS is managing resources \*Resource allocation policies must consider. Different Architectural Approaches: Demands on operating systems require new ways of organizing the OS.  
MicroKernel Architecture: \*assigns only a few essential functions to the kernel:  
Multithreading: \*technique in which a process, executing an application, is divided into threads that can run concurrently.  
Symmetric Multiprocessing (SMP): \*term that refers to a computer hardware arch. and also to the OS behavior that exploits that arch. \*several processes can run in parallel. \*multiple processors are transparent to the user. \*these processors share the same main memory and I/O facilities. \*all processors can perform the same functions. \*The OS takes care of scheduling of threads or processes on individual processors and of synchronization among processors.  
OS Design: [Distributed OS] \*provides the illusion of 1) single main memory space 2) single secondary memory space 3) unified access facilities. \*State of the art for distributed operating systems lags that of uniprocessor and SMP operating systems.  
[Object-oriented Design] \*used for adding modular extensions to a small kernel. \*enables programmers to customize an operating system without disrupting system integrity. \*eases the development of distributed toos and full-blown distributed operating systems.  
Virtual Machine and Virtualization: enables a PC or server to simultaneously run multiple operating systems or multiple sessions of a single OS. \*a machine can host numerous application, including those that run on different operating systems, on a single platform. \*host operating system can support a number virtual machines.  
\*each has the characteristics of a particular and, in some versions of virtualization, characteristics of a particular hardware.



OS Management of Application Execution: \*resources made avail. to mul. appli. \*The proc is switched among mul. appli. so all will appear to be progressing. \*the proc and I/O devices can be used efficiently.  
Process Control block: \*contains the process elements. \*it is possible to interrupt a running process and later resume exe as if the interrupt had not occurred. \*key tool that allows support for multiple proc.

Two state process model: state 1) running state 2) not running.  
Reasons for Process creation: 1) new batch job 2) interactive logon 3) created by OS to provide a service 4) spawned by existing process.  
Process termination: \*there must be a means for a process to indicate its completion. \*a batch job should include HALT instruction or an explicit OS service call for termination. \*for an interactive application, the action of the user will indicate when the process is completed (e.g. logoff, quitting an application).  
Reasons for process termination: 1) normal completion 2) time limit exceeding 3) memory unav 4) bounds viol 5) protection err 6) arithmetic err 7) I/O failure 8) parent termination  
Suspended Processes: \*Swapping - involves moving part of all of a process from main memory to disk \*when none of the processes in main memory is in the ready state, the OS swaps on of the blocked process out on to disk into a suspend queue.  
Characteristics of a Suspended Process: \*the process is not immediately available for exe. \*The process may or may not be waiting on an event. \*The process was placed in a suspended state by an agent: either itself, a parent process, or the OS, for the purpose of preventing its exe. \*the process may not be removed from this state until the agent explicitly orders the removal.  
Reasons for suspended process: 1) swapping, 2) Os may suspend a background or utility 3) interactive user request 4) timing, may be executed periodically 5) parent process request, parent may wish to suspend exe of a descendent to modify.  
OS CONTROL TABLES  
Memory Tables: \*used to keep track of both main and secondary memory. \*Processes are maintained on secondary memory using some sort of virtual memory or simple swapping mechanism. Must include \*allocation of main mem to process \*allocation secondary mem to process \*protection attributes of blocks of main or virtual mem. \*information needed to manage virtual mem.  
I/O tables: \*used by the OS to manage the I/O devices and channels of the computer system. \*at any given time, an I/O device may be available or assigned to a particular process.  
File Tables: \*information may be maintained and used by a file management system in which case the OS has little or no knowledge of files. \*In other OS, much of the detail of file mangement is managed by the OS itself. \*existence of files \*location of secondary mem \*current status  
Process Tables: Must be maintained to manage processes. \*there must be some reference to memory, I/O, and files, directly or indirectly. \*The tables themselves must be accessible by the OS and therefore are subject to memory management.  
PROCESS CONTROL STRUCTURE \*OS must know \*where the process is located \*The attributes of the process that are necessary for its management.  
Process Control Location: Process location - \*Process must include a program to be exe, \*process will consist of at least sufficient memory to hold the programs and data of the process. \*the execution of a program typically involves a stack that is used to keep track of procedure calls and parameter passing between procedures.  
Process Attributes: each process has associated with it a number of attributes controlled by the OS. \*The collection of program, data, stack, attributes is referred to as process image. (location depends on memory management scheme being used.)  
Process Image: 1) user data (modifiable part of user space, stack, etc.) 2) User program 3) stack (each process has one or more LIFO stacks associated. stack is used to store parameters and calling addresses for system calls.) 4) Process control block (data needed by the OS to control the process.  
Process Identification: \*each process is assigned a unique numeric ID, or there must be a mapping that allows the OS to locate the appropriate tables based on the process ID. \*many of the tables controlled by the OS may use process ID to cross-reference process tables. \*Memory tables may be organized to provide a map of main memory. \*When processes communicate with each other the process ID informs the OS of the destination of a particular communication. \*when processes are allowed to create other processes, identifiers indicate the parent and descendent.  
Processor State information: Consists of the contents of the process registers \*user-visible registers \*control and status registers \*stack pointers  
Program Status Word (PSW): contains condition codes plus other status information. \*EFLAGS register is an example of a psw used by any OS running on an x86 proc.  
Process Control Information: the additional info needed by the OS to control and coordinate the various active processes.  
Process Control Block: \*most important data structure in an OS \*contains all of the information about a process that is needed by the OS. \*blocks are read and/or modified by virally every module in the OS. \*Defines the state of the OS. \*Difficulty is not access but protection. \*a bug in a single routine could damage process control blocks, which could destroy the system's ability to manage the affected process. \*A design change in the structure or semantics of the Process control block could affect a number of modules in the OS.

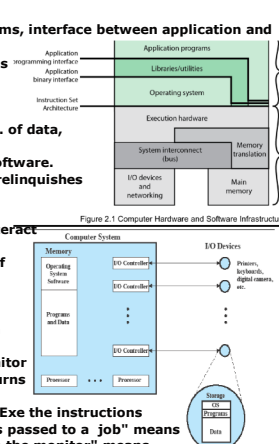


Figure 2.1 Computer Hardware and Software Infrastructure




Figure 2.2 The Operating System as Resource Manager

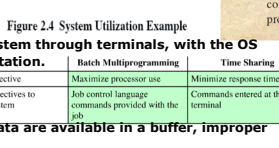


Figure 2.4 System Utilization Example

Processor is often idle  
even with automatic job sequencing  
I/O devices are slow compared to processor  
Read one record from file 15 μs  
Execute 100 instructions 1 μs  
Write one record to file 15 μs  
TOTAL 31 μs  
Percent CPU Utilization =  $\frac{1}{31} = 0.032 = 3.2\%$   
Principal objective Batch Multiprogramming Time Sharing  
Maximize processor use Minimize response time  
Job control language commands provided with the job Commands entered at the terminal  
Source of directives to operating system  
Figure 2.10 Virtual Memory Addressing